





Microservices latency troubleshooting guide

IPPON en quelques chiffres



15 years
2002 - 2017



325
Consultants



31 M€
2017 revenue



40 M€
2018 forecast



4
continents

Plan

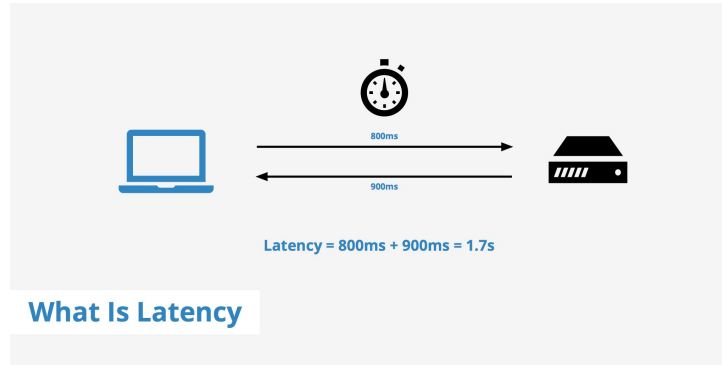
1. Why should I care ?
2. Do I really need tracing ?
3. Tracing basics
4. What should I use ?
5. Demo time

A black and white photograph of a swimmer in a pool, captured in the middle of a stroke. The swimmer's head is above water, and their arms are extended forward, creating a large, dynamic splash of water. The background is dark, making the white water splash stand out. The text "Why should I care" is overlaid on the left side of the image.

Why should I care

The latency

- *"Latency is a time interval between the stimulation and response, or, from a more general point of view, a time delay between the cause and the effect of some physical change in the system being observed."*



keycdn.com

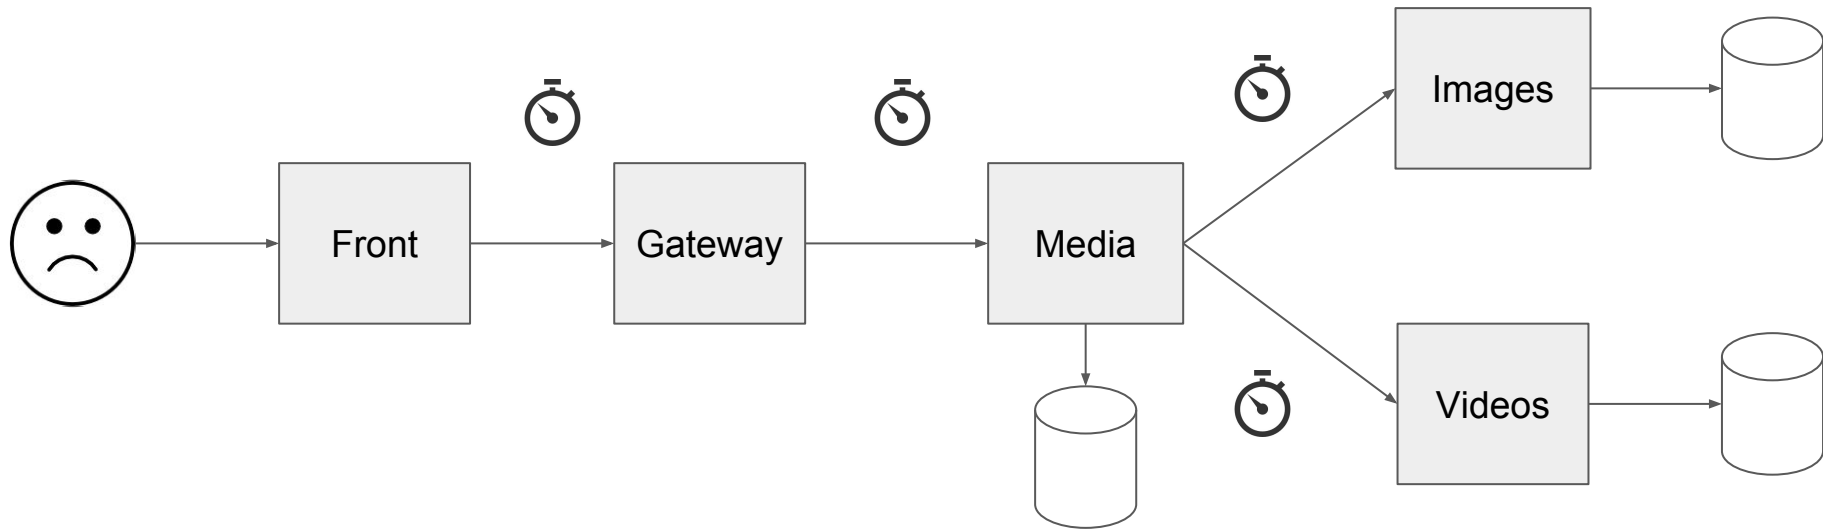
- It has impact
 - Amazon found every 100ms of latency cost them 1% in sales.
 - Google found an extra 0.5 seconds in search page generation time dropped traffic by 20%.
 - A broker could lose \$4 million in revenues per millisecond if their electronic trading platform is 5 milliseconds behind the competition

Enters the microservices

- Has some famous examples
 - Netflix
 - Google
 - AirBnB
- Fits well with the way we manage infrastructure
 - (multi) Cloud
 - Containers
 - FAAS
- Brings a lot of new challenges
 - Deployment
 - Interservice communication
 - **Monitoring**




Latency and network




What if I use a queue ?


- Changes nothing in term of latency
- Latency is still a key indicator
- You still have to use network to send messages
- This is just another transport mechanism !




Simple




Basic Concurrency



Async Concurrency



Distributed Concurrency



And there is more

- You need to align clock between different instances
 - Use timestamps ?
 - Or to treat time at collect time ?
- How to gather data without impacting performances ?
- Does It work with thousands of services ?
- This is not trivial





Do I really need tracing ?

Can't I just use logs ?

- *"In computing, a log file is a file that **records** [...] **events** that occur in an operating system or other software runs [...] Logging is the act of keeping a log."*
- We are at the **event scale**
- The good part
 - Simple timestamped and structured messages
 - Tools are available to do it, easily
 - Easy to grep, and read manually
 - Easy to aggregate and parse
- Overhead ?
 - Grows with traffic and verbosity
 - Storage impact limited by pipeline filtering
- Perfect for monoliths, black boxes, exceptional cases

What about metrics?

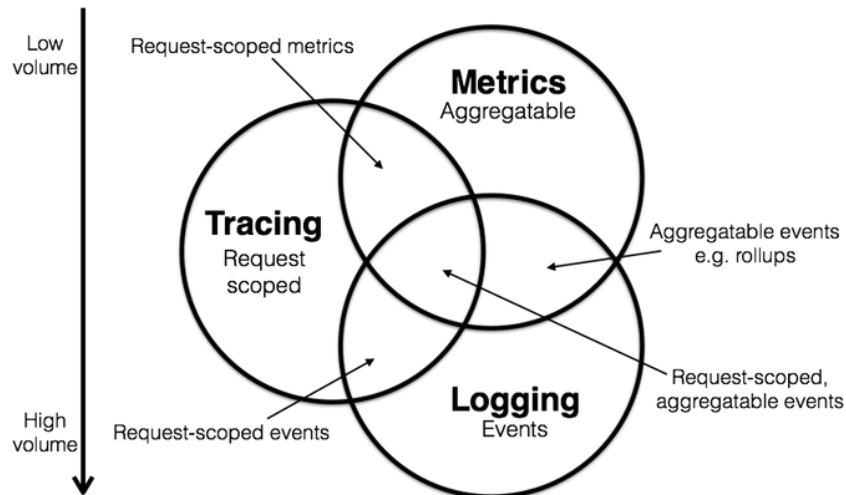
- *“A software metric is a standard of measure of a degree to which a software system or process possesses some property.”*
- ≈ Data combined from measuring events
- We are at the **software scale**
- The good part
 - Time and store values of anything
 - Tools are available to do it
 - Report duration buckets near-real time
 - Let you identify trends
- Overhead
 - Growing size is fixed
 - Storage impact limited by writing only what you need
- Perfect for identifying patterns and for alerting

What is tracing?

- *"In software engineering, tracing involves a specialized use of logging to record information about a program's execution [...]. **Tracing is a cross-cutting concern.**"*
- ≈ Recording events with causal ordering
- We are at the **system scale**
- The good part
 - Common terminology
 - Tools are available to do it
 - Report data near-real time
 - Identify issues across services
- Overhead
 - Grows with traffic
 - Storage impact limited by various techniques
- Perfect for distributed services


What you want ?

- Use everything
- They work together
 - Log correlation with traces
 - Correlate traces with metrics
 - Fetch metrics on collected traces
- Production issue debugging
 - Something is wrong
 - Lookup metrics for abnormal load, broken instances
 - Check for traces with abnormal latency, errors
 - Find the impacted service
 - Lookup for the logs in a particular service for a particular trace ID
 - Fix the issue
 - Enjoy your night



Source: Peter Bourgon

Tracing basics

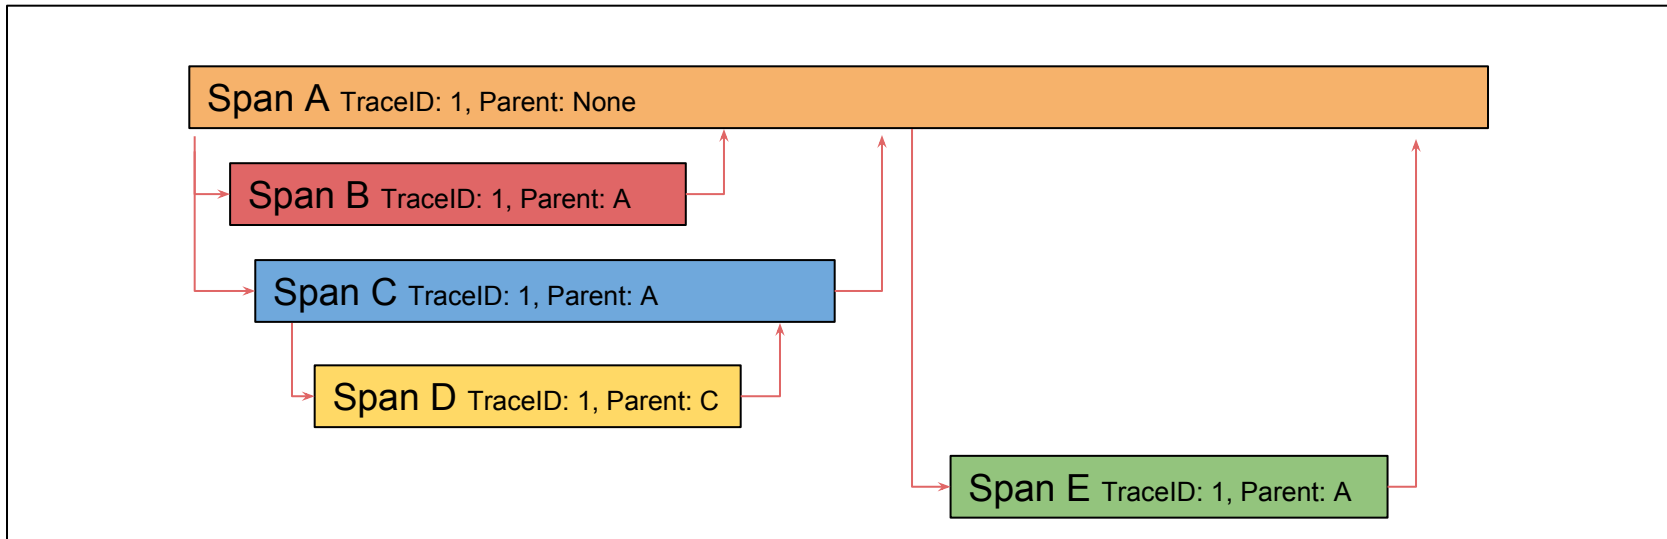


A bit of history

- Google used the dapper paper for more than ten years
- Been published online in 2010 by a Google team
- Twitter open sources Zipkin in 2012
- OpenTracing created in 2014
- OpenTracing joins the CNCF in 2016
- [First Draft of trace context published](#) by a W3C working group in 2018

A **Trace** represent your whole request across all services involved.

A **Span** is the basic unit of work in a given Trace. You can see it as the duration of an operation.



service2.http:/readtimeout: 3.557s



AKA: service1,service2

| Date Time | Relative Time | Annotation | Address |
|----------------------|---------------|----------------|---------------------------|
| 19/12/2016, 14:19:23 | 307.000ms | Client Send | 127.0.0.1:8081 (service1) |
| 19/12/2016, 14:19:23 | 310.000ms | Server Receive | 127.0.0.1:8082 (service2) |
| 19/12/2016, 14:19:26 | 3.836s | Server Send | 127.0.0.1:8082 (service2) |
| 19/12/2016, 14:19:27 | 3.864s | Client Receive | 127.0.0.1:8081 (service1) |

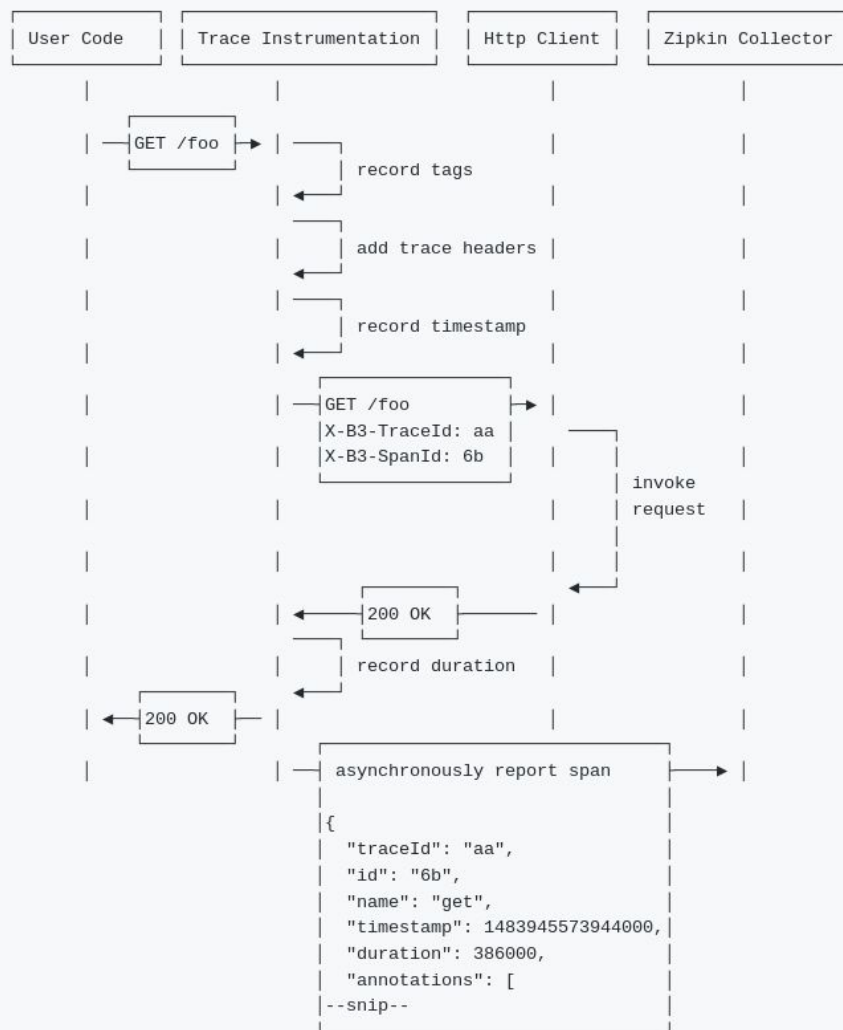
| Key | Value |
|-----------------------|---|
| error | Request processing failed; nested exception is org.springframework.web.client.ResourceAccessException: I/O error on GET request for "http://localhost:8082/blowup": Read timed out; nested exception is java.net.SocketTimeoutException: Read timed out |
| http.host | localhost |
| http.method | GET |
| http.path | /readtimeout |
| http.status_code | 500 |
| http.url | http://localhost:8082/readtimeout |
| mvc.controller.class | BasicErrorController |
| mvc.controller.method | error |

- Tags

- Key value pairs
- Additional data on spans
- Not shared with other spans
- Ex : http.host, http.method ...

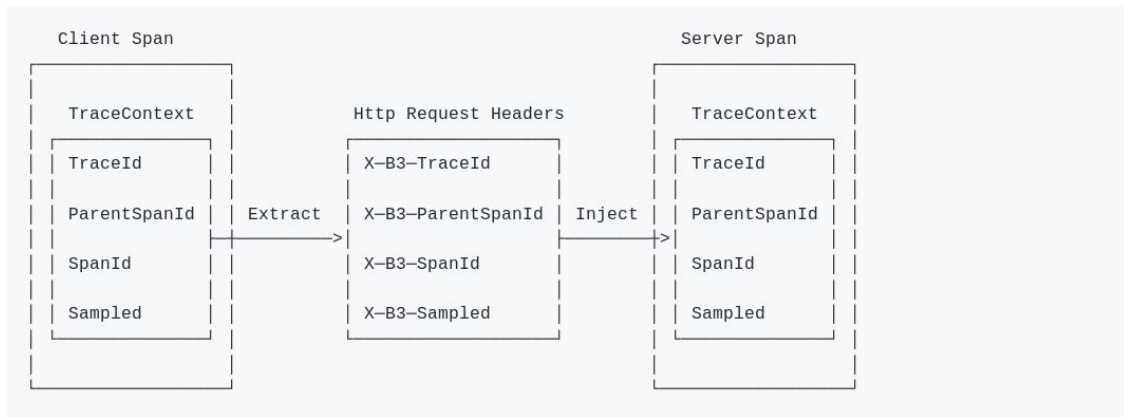
- Baggages

- Key value pairs
- Part of the propagated context
- Be careful about latency impact if you add too many



How to keep context

- Inbound / Outbound
 - Read/Write request headers
 - Read/Write queue headers
 - Read/Write request body



- We need to keep a context per request
 - ThreadLocal
 - Pass context between methods

Visualisation

Duration: 167.142ms Services: 6 Depth: 15 Total Spans: 65

Expand All Collapse All Filter Ser... JSON

brewing x35 broker x16 ingredients x7 presenting x8 reporting x16 zuul x9



- On the wire the cost is low
 - Few headers
 - Baggage
- Span reporting
 - batch
 - asynchronous
- Sampling
 - Percentage
 - User defined
- From a developer point of view
 - Transparent most of usual cases
 - Need awareness for advanced cases

A grayscale photograph of two football players from behind, standing on a field. The player on the left wears jersey number 86, and the player on the right wears jersey number 17. Both are wearing helmets and jerseys. The background is a blurred field with spectators.


What should I use ?

Standardization ?

- CNCF and Opentracing
 - More or less an interface to implement to be compliant
 - Is a good idea
 - Lot of [implementations](#)
 - But
 - wire format not fixed
 - Will create a lot of tracing library, maybe incompatible
 - Maybe too much responsibility
- [W3C trace context headers](#) working group



- Zipkin
 - One of the historic tracing system
 - Nearly all languages supported
 - Initiated by Twitter, now community driven
- Jaeger
 - One of Opentracing first implementation
 - Has a lot of traction
 - Backed by Uber, now in the CNCF
- Opencensus
 - The newcomer
 - Alpha stage
 - Backed by [Google & Microsoft](#)



In the cloud



TRACE REQUESTS

AWS X-Ray traces requests made to your application.

X-Ray collects data about the request from each of the underlying application services it passes through.

RECORD TRACES

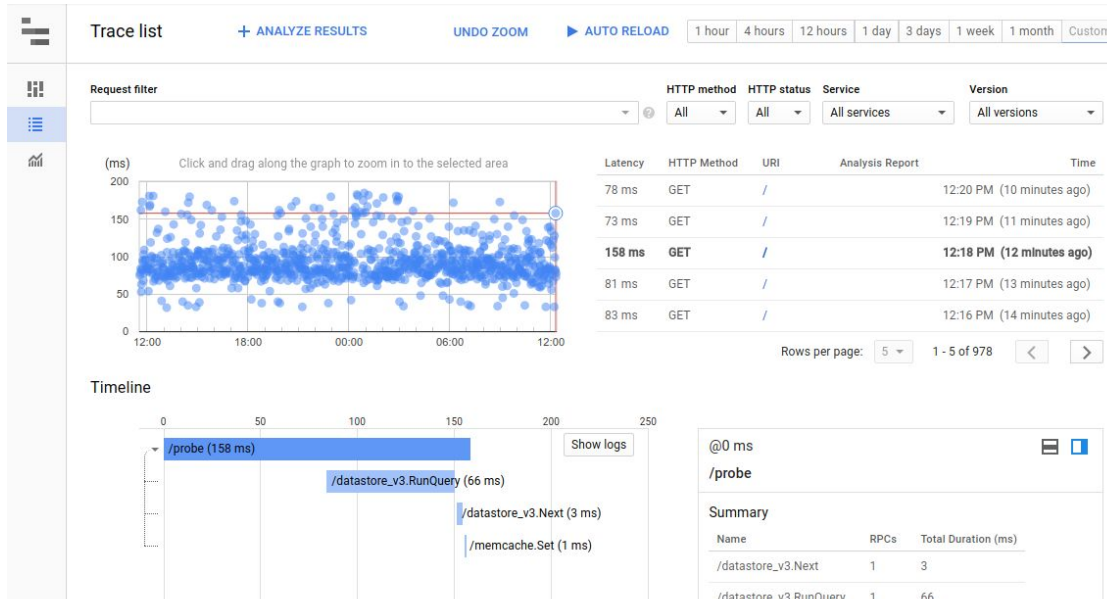
X-Ray combines the data gathered from each service into singular units called traces.

VIEW SERVICE MAP

View the service map to see trace data such as latencies, HTTP statuses, and metadata for each service.

ANALYZE ISSUES


Drill into the service showing unusual behavior to identify the root issue.



Without code

- Services mesh add new monitoring capabilities
 - Automatically collect requests
 - The apps needs to propagate headers

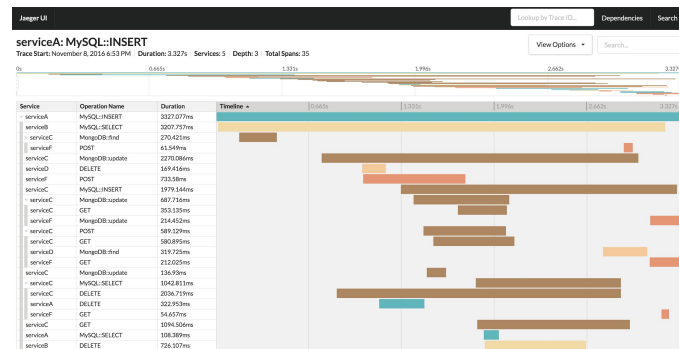
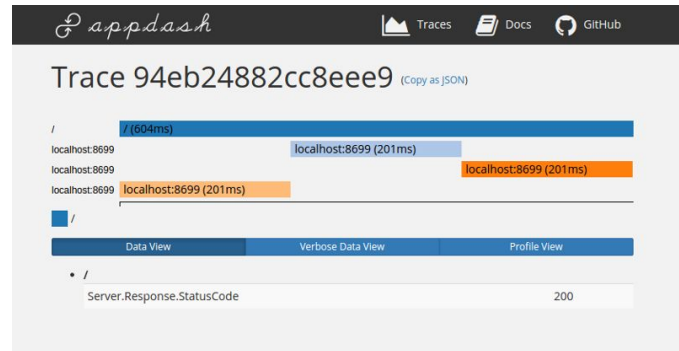
- APM
 - Instana
 - Skywalking
 - ...



Visualisation tools

- Zipkin
- Lightstep
- Jaeger
- appdash
- Tracer
- Stackdriver
- X-ray
- ...

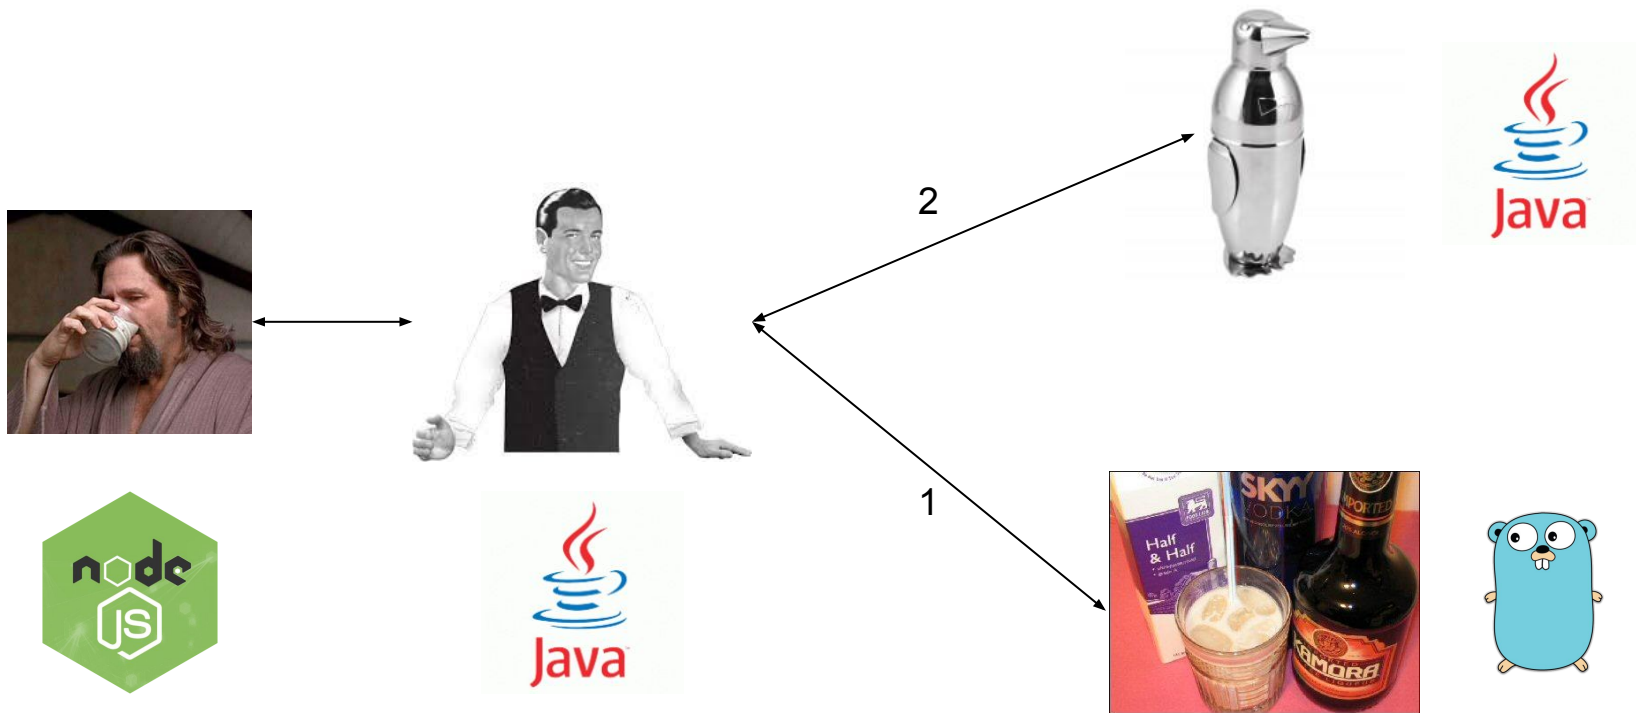
- Backed by a storage infrastructure
 - Elasticsearch
 - Cassandra
 - MySQL
 - Cloud magic



A dark, grainy background image showing the silhouette of a person standing on a surfboard, riding a wave. The person is in a dynamic pose, leaning forward with arms outstretched. The surfboard is visible beneath their feet, and the wave's crest is breaking behind them.

Demo time

The example



```
private Tracer tracer;

private void getGlass() throws InterruptedException {
    Span s = tracer.nextSpan().start();
    s.name("glass");
    logger.info("Fetching a glass");
    s.tag("type", "old fashioned");
    Thread.sleep(50);
    s.annotate("fetched");
    Thread.sleep(100);
    s.annotate("cleaned");
    s.finish();
}
```

NodeJS integration

```
1 // Continuation local storage
2 var ctxImpl = new CLSContext('zipkin');
3
4 // Add a http transport to local zipkin instance
5 var recorder = new BatchRecorder({
6   logger: new HttpLogger({
7     endpoint: 'http://localhost:9411/api/v1/spans'
8   })
9 });
10
11 // Create a tracer
12 var tracer = new Tracer({ctxImpl, recorder});
13
14 // Add the Zipkin middleware to express js
15 app.use(zipkinMiddleware({
16   tracer,
17   serviceName: 'dude', // name of this application
18   sampler: new zipkin.sampler.CountingSampler(1)
19 }));
20
21 // instrument the client
22 var zipkinRest = rest.wrap(restInterceptor, {tracer, serviceName: 'dude'});
23
24 // Configure REST endpoints
25 app.get('/cocktail', (req,res) => {
26   zipkinRest('http://localhost:8081/make')
27     .then(
28       (response) => res.send(response.entity),
29       (response) => console.error("Error", response.status)
30     )
31 });
```

Go integration

```
// Initialize zipkin http collector
collector, err := zipkintracer.NewHTTPCollector("http://localhost:9411/api/v1/spans")
// Create the tracer
tracer, err := zipkintracer.NewTracer(
    zipkintracer.NewRecorder(collector, false, ":8082", "clerk"),
)
opentracing.InitGlobalTracer(tracer)
// Define handlers
http.HandleFunc("/fetchIngredients", func(w http.ResponseWriter, r *http.Request) {
    // Fetch the request context
    wireContext, err := opentracing.GlobalTracer().Extract(
        opentracing.HTTPHeaders,
        opentracing.HTTPHeadersCarrier(r.Header))
    if err != nil {
        log.Print(err)
    }
    // Create a new span from the context
    serverSpan := opentracing.StartSpan("clerk", ext.RPCServerOption(wireContext))
    defer serverSpan.Finish()

    // We need to explicitly pass context to other methods
    getMilk(serverSpan.Context())
    fmt.Fprint(w, "Good milk and ice for you !")
})
```

Wrapping up



- You need distributed tracing to
 - Do root cause analysis
 - Understand your system
 - Optimize performances
 - Increase efficiency
 - Improve reliability
- Collaborate wells with other monitoring tools
 - Log
 - Metrics
 - Traces
- Various tools allow you to do this easily
 - With tracers
 - With framework instrumentation
 - With cloud first tools
 - With your favourite APM
- This is a moving field
 - Standards are emerging
 - New tracers & instrumentations
 - A must have for distributed systems !

Useful resources

- Code

- [My demo](#)
- [Brewing demo](#)
- [OpenZipkin project](#)
- [Jaeger tracing project](#)

- Slides & conferences

- [FullstackFest conference](#) by @adrianfcole
- [MicroXchg conference](#) by @jcchavez
- [Kubecon conference](#) by @YuriShkuro
- [Slides from @el_bhs](#)

- Blogs

- [A blog on observability](#)
- [Metrics Tracing & Logging by @peterbourgon](#)
- [Distributed Tracing core concepts by @MunroeNic](#)



Digital . Technologies . Hosting

PARIS
BORDEAUX
NANTES
LYON
MARRAKECH
WASHINGTON DC
NEW-YORK
RICHMOND
MELBOURNE

contact@ippon.fr
www.ippon.fr - www.ippon-hosting.com - www.ippon-digital.fr
@ippontech

-
01 46 12 48 48